

The operator antipattern

Kubernetes Community Days London 2024

Gerald Schmidt





Source: simpsons.fandom.com/wiki/Springfield_Box_Factory

Cindy Sridharan

@copyconstruct · [Follow](#)



As more and more exec types begin to wonder why companies are overstaffed, we're going to enter an era where engineers are going to be questioned why they run such complex systems - think every tech trend of the 2010's - when much simpler ones would suffice.

Brace yourselves.

4:19 PM · Nov 7, 2022



598



Reply

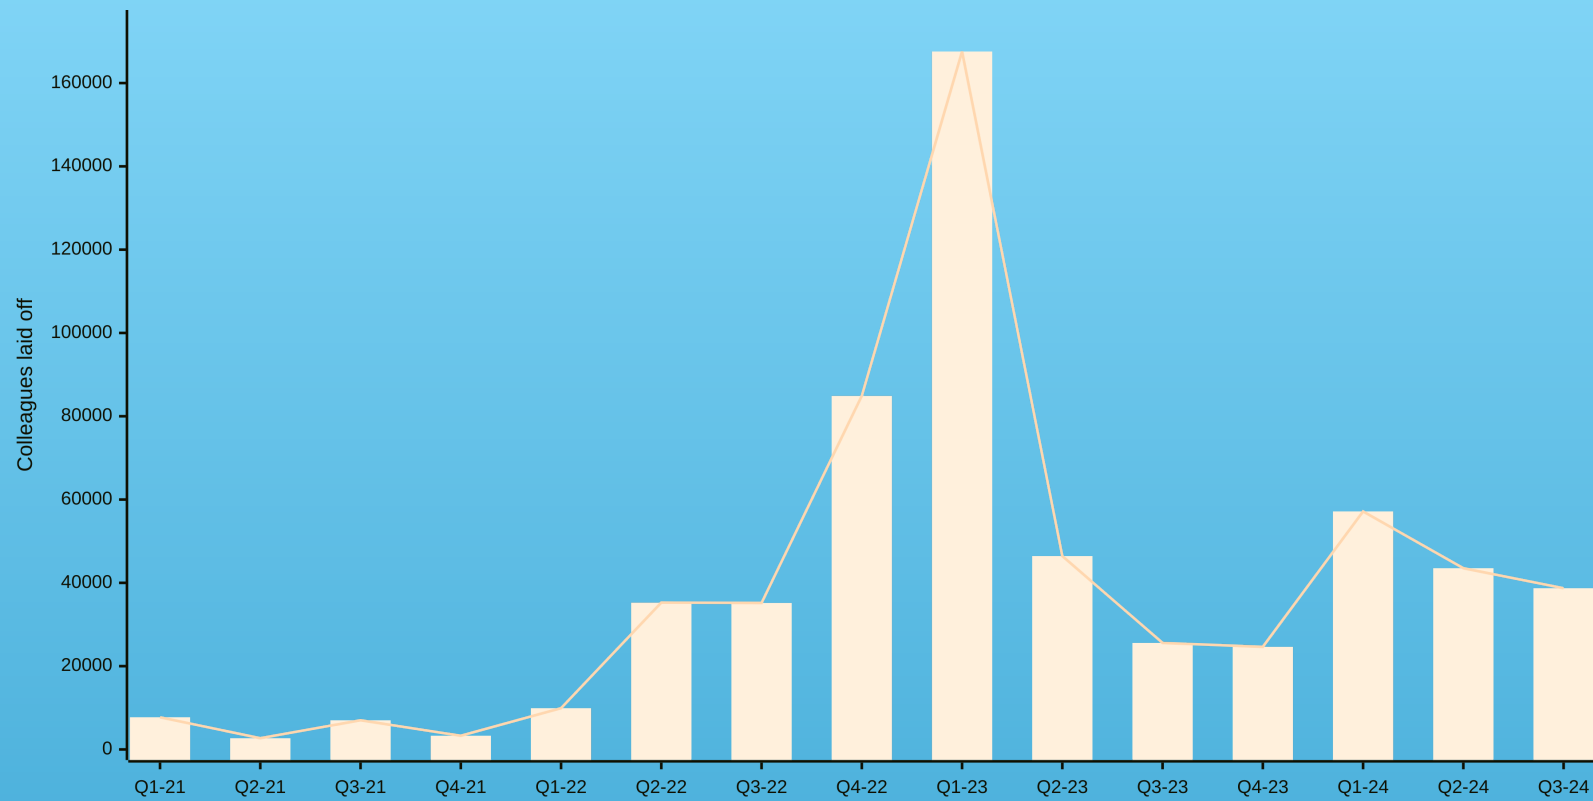


Copy link

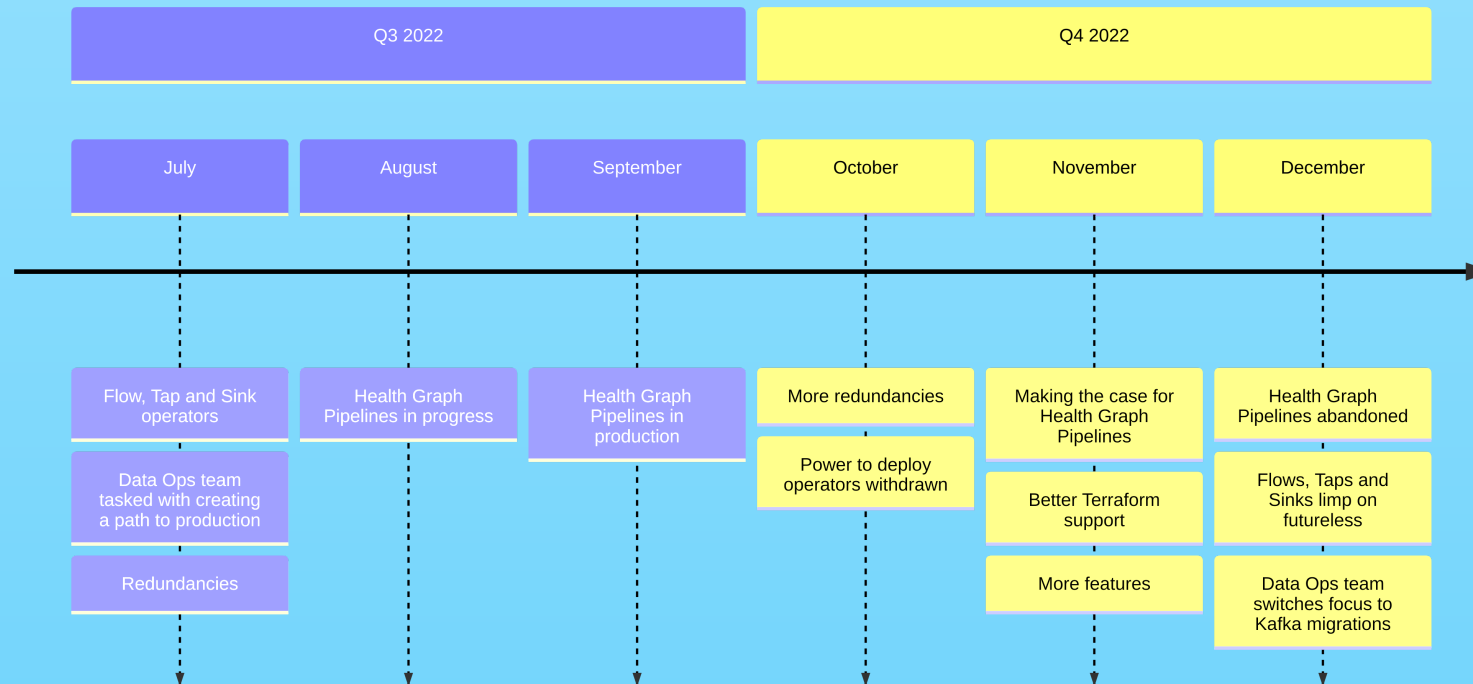
[Read 17 replies](#)

Classic Sridharan:

[Testing Microservices, the sane way \(2017\)](#) | [Testing in Production, the safe way \(2018\)](#) | [Testing in Production: the hard parts \(2019\)](#)



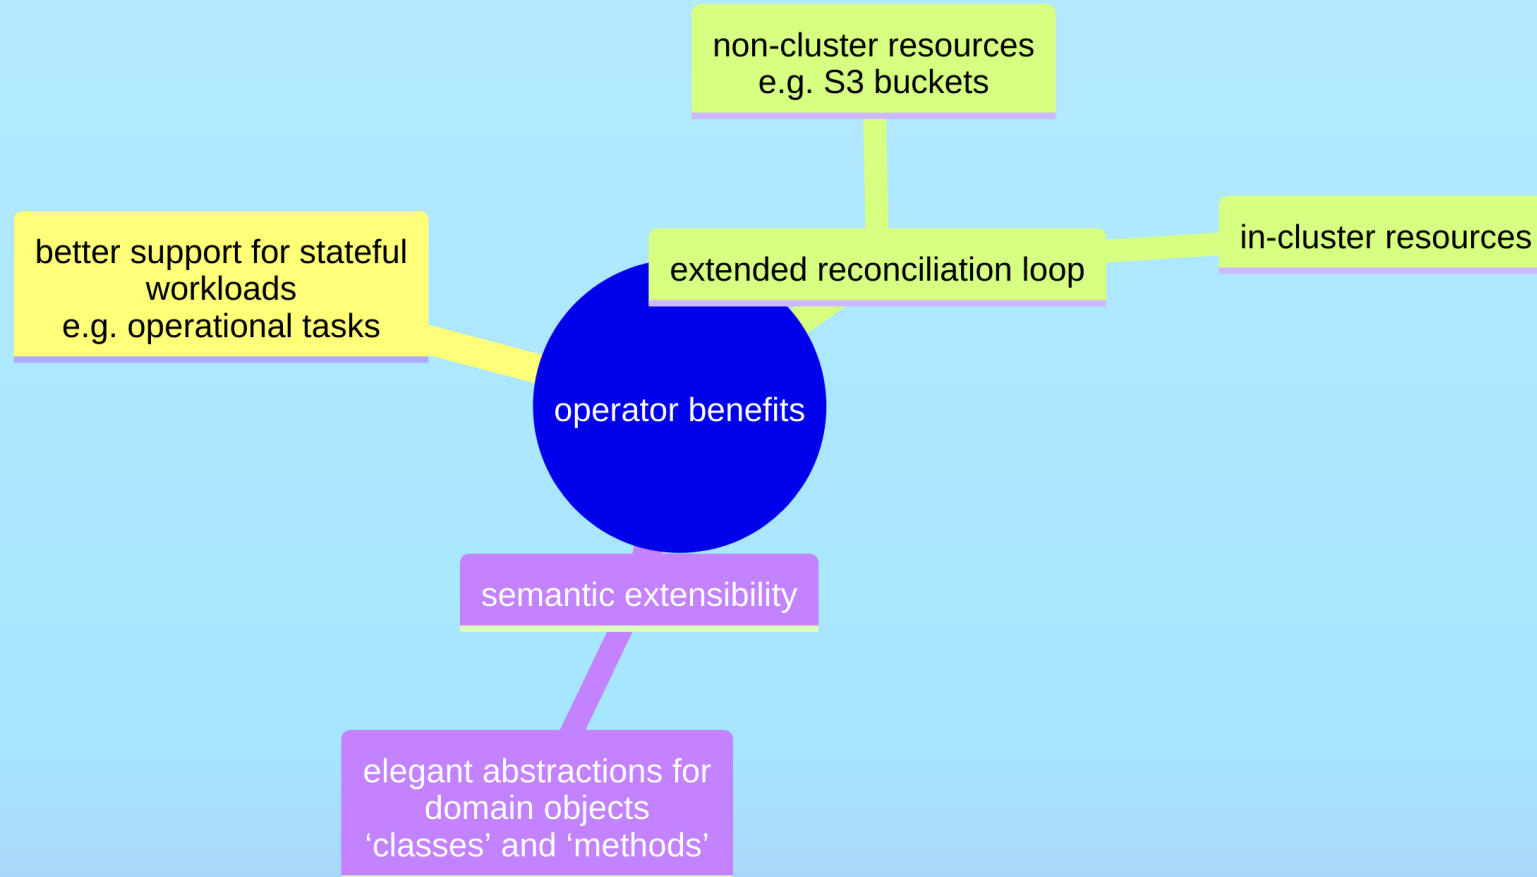
Source: layoffs.fyi



The premise

An Operator is an application-specific controller that extends the Kubernetes API to create, configure, and manage instances of complex stateful applications on behalf of a Kubernetes user. It builds upon the basic Kubernetes resource and controller concepts but includes domain or application-specific knowledge to automate common tasks.

Brandon Philips, [Introducing Operators: Putting Operational Knowledge into Software](#) (2016)



The CRD/controller split


Weighting

The benefits of operators centre on the controller.

The custom resource component is to a meaningful extent syntactic sugar.

At the same time it is responsible for most of the problematic side-effects of operators: custom resource definitions are cluster-level objects, which brings with it a whole set of permission issues. Versioning is another problematic aspect that we'll come back to.

Cost and Results of experiments



Prometheus without operator

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   annotations:
5     prometheus.io/port: "2112"
6     prometheus.io/scrape: "true"
```


Prometheus with operator

```
1 $ kubectl get crd -o custom-columns=NAME:.metadata.name \  
2   | grep "monitoring\.coreos\  
3 alertmanagerconfigs.monitoring.coreos.com  
4 alertmanagers.monitoring.coreos.com  
5 podmonitors.monitoring.coreos.com  
6 probes.monitoring.coreos.com  
7 prometheusagents.monitoring.coreos.com  
8 prometheuses.monitoring.coreos.com  
9 prometheusrules.monitoring.coreos.com  
10 scrapeconfigs.monitoring.coreos.com  
11 servicemonitors.monitoring.coreos.com  
12 thanosrulers.monitoring.coreos.com
```

Kube Prometheus stack

```
1 FILE=charts/kube-prometheus-stack/README.md
2 for TAG in "9.3.4" "15.0.0" "60.0.0"; do
3     git checkout "kube-prometheus-stack-${TAG}" 2>/dev/null
4     echo "# ${TAG}"
5     head -n -100 "${FILE}" | wc -l charts/kube-prometheus-stack/README.md
6     grep -v "\(^kubect1 .*crd\|CRD\) " "${FILE}" | wc -l
7 done
```

Kube-prometheus-operator



Antipattern 1: operators in developer workflows

Flow operator

Antipattern 2: tight coupling with external resources

Strimzi

Antipattern 3: versioning trouble

Incrementing CRD versions is a serious matter.

Is the old version still served? Have we provided a conversion webhook?

So far from reducing complexity, we are introducing new error conditions, failure modes and edge cases.

See AWS Controllers for Kubernetes. (Still on alpha.)

Antipattern 4: overpromising

The association of operators with complex *stateful* applications has not displaced managed databases such as the Relational Database Service. The CRD that allows me to create a VectorDatabase resource does not magically make it a good, fault-tolerant. A backup method is helpful and appreciated, but it does not rival a mature point-in-time recovery facility.

Whisper it: Kubernetes does not have a 'stateful workload' problem. It has a persistent volume problem. The solution is object storage, and the challenge is working around the limitations of object storage when it comes to read and write speed.

[See Object storage for stateful applications on Kubernetes](#)

Kyverno

Wins first prize for an implementation that feels as if it should be an in-tree policy engine. Policy violations create detailed events and the new resources (Policy, ClusterPolicy) fit well into the existing set of resources.

Controller revival

Grafana has bucked the trend of CRD sprawl.


To load a dashboard on startup, Grafana seeks out ConfigMaps that have label `grafana_dashboard` set to value `1`.

There is no need for a GrafanaDashboard CRD.

Grafana dashboards are JSON objects following a well-established structure.

Teams store dashboards they wish to keep in a folder:

```
1 for DASHBOARD in \
2   $(ls kube-prometheus-stack/dashboards/*.json)
3 do
4   CONFIGMAP=$(basename "${DASHBOARD}" | cut -d'.' -f1)
5   kubectl create configmap "${CONFIGMAP}" \
6     -n monitoring \
7     --dry-run=client \
8     --from-file="${DASHBOARD}" -o yaml | \
9     kubectl apply -f -
10  kubectl label configmap "${CONFIGMAP}" \
11    -n monitoring \
12    --overwrite grafana_dashboard="1"
13 done
```

 [gerald1248/operator-antipattern-slides](#)
 www.linkedin.com/in/gerald1248
 [03spirit](#)

Slides built with [Markdeck](#)

